



Building an Event Driven Gen-AI powered video summarizer API



Bonjour !

I am Zied Ben Tahar

AWS Community Builder / Principal Architect @ Aviv Group /

You can find me:

Blog: <https://zied-ben-tahar.medium.com/>

Linkedin: <https://www.linkedin.com/in/zied-ben-tahar-5200913/>

Twitter : [@wow_sig](https://twitter.com/wow_sig)

Agenda

01

GenAI on AWS - A
quick overview

02

Let's build: AI powered
Video summarizer

03

Video summarizer -
Scaling the
architecture with EDA

04

Api-fying the solution

05

Demo Time

A large red circle is positioned at the top center. Several thin red lines and smaller red circles are scattered across the right side of the slide, creating a dynamic, abstract background.

01

GenAI on AWS - A quick overview



Amazon Bedrock

Amazon Bedrock

A serverless Gen-AI service

Managed FMs

Access to foundation models
from leading partners on AI space

Customizable

Capabilities to adapt and
privately customize FMs to
domain specific problem

Serverless

No infrastructure to manage,
Good integration with AWS
ecosystem



AI21 Labs Jurassic



Amazon Titan



Anthropic's Claude



Cohere Command



Meta Llama 2



Stability AI SDXL



MISTRAL
AI_

Gen AI

Addressing some LLM challenges

How to solve issues related to knowledge cut-off, hallucinations, blackbox ?

→ **Prompt engineering**

- ◆ Optimizing prompts for specific tasks and guiding the model to yield the expected answers

→ **Retrieval augmented generation (RAG)**

- ◆ Allow the model to access to external/up to date information to enhance response.

→ **Model fine-tuning**

- ◆ Additional training of a pre-trained model on task-specific datasets.

Anthropic's Claude ~~(2.1)~~ 3

A serious alternative to GPT 4

- Handles large context windows, 200K tokens (150k words)
- Able to maintain factual consistency when processing long prompts
- Can generate code and structured data
- Vision capabilities (new with Sonnet)

	Claude 3 Opus	Claude 3 Sonnet	Claude 3 Haiku	GPT-4	GPT-3.5	Gemini 1.0 Ultra	Gemini 1.0 Pro
Undergraduate level knowledge <i>MMLU</i>	86.8% 8-shot	79.0% 8-shot	75.2% 8-shot	86.4% 8-shot	70.0% 8-shot	83.7% 8-shot	71.8% 8-shot
Graduate level reasoning <i>GPQA, Diamond</i>	50.4% 0-shot CoT	40.4% 0-shot CoT	33.3% 0-shot CoT	35.7% 0-shot CoT	28.1% 0-shot CoT	—	—
Grade school math <i>GSMSK</i>	95.0% 0-shot CoT	92.3% 0-shot CoT	88.9% 0-shot CoT	92.0% 5-shot CoT	57.1% 5-shot	94.4% Maj1@32	86.5% Maj1@32
Math problem-solving <i>MATH</i>	60.1% 0-shot CoT	43.1% 0-shot CoT	38.9% 0-shot CoT	52.9% 4-shot	34.1% 4-shot	53.2% 4-shot	32.6% 4-shot
Multilingual math <i>MGSM</i>	90.7% 0-shot	83.5% 0-shot	75.1% 0-shot	74.5% 8-shot	—	79.0% 8-shot	63.5% 8-shot
Code <i>HumanEval</i>	84.9% 0-shot	73.0% 0-shot	75.9% 0-shot	67.0% 0-shot	48.1% 0-shot	74.4% 0-shot	67.7% 0-shot
Reasoning over text <i>DROP, F1 score</i>	83.1 3-shot	78.9 3-shot	78.4 3-shot	80.9 3-shot	64.1 3-shot	82.4 Variable shots	74.1 Variable shots
Mixed evaluations <i>BIG-Bench-Hard</i>	86.8% 3-shot CoT	82.9% 3-shot CoT	73.7% 3-shot CoT	83.1% 3-shot CoT	66.6% 3-shot CoT	83.6% 3-shot CoT	75.0% 3-shot CoT

Claude 3 benchmarks



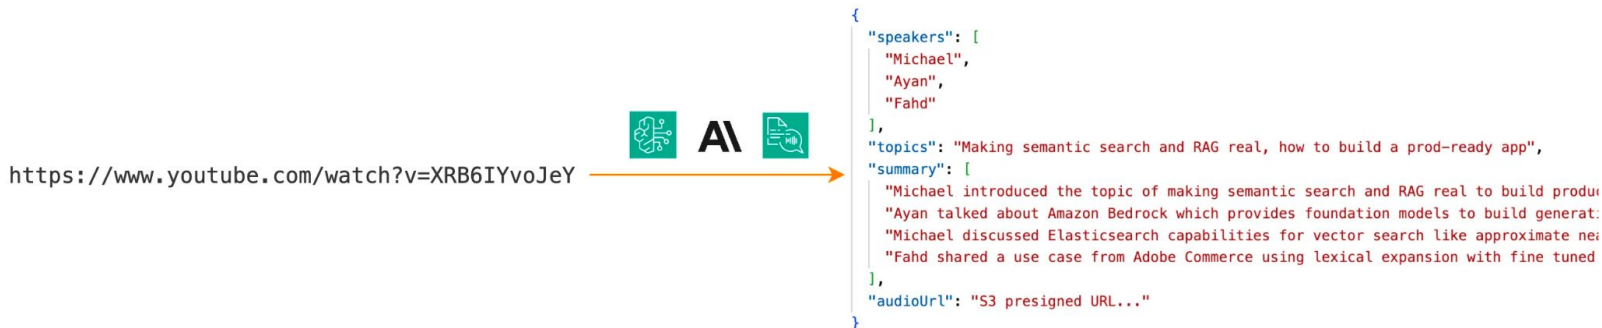
02

Let's build!
AI powered Video
summarizer



Solution overview

Summarization workflow, first attempt



01

Extract/Get Video
transcript

02

Summarize & extract
structured data

03

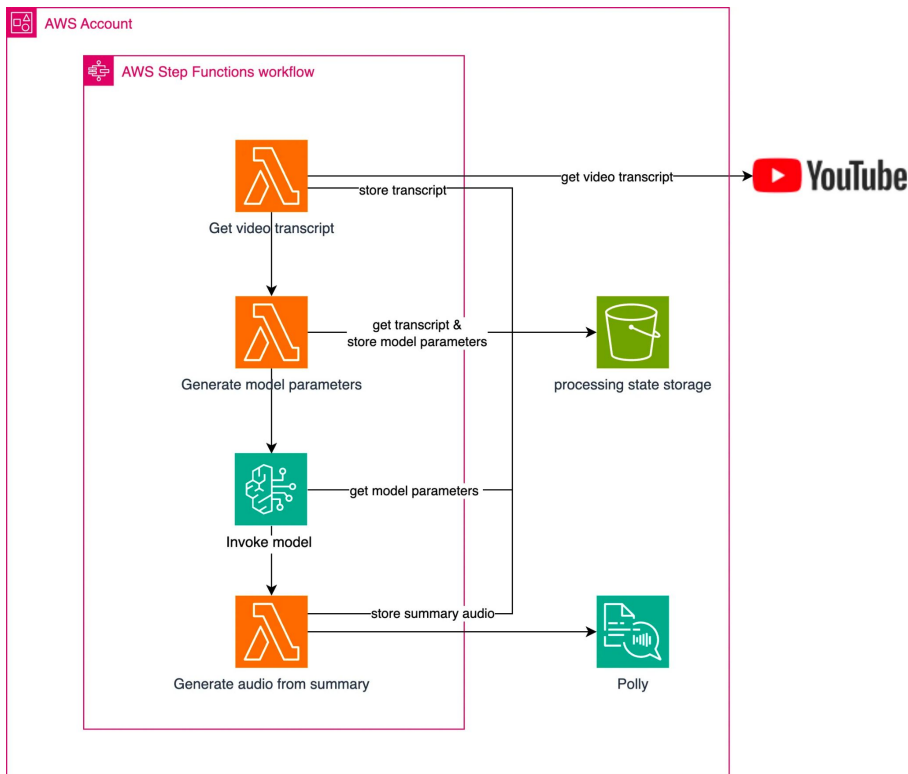
Generate audio from
summary

04

Profit !

Solution overview


Serverless Summarizer workflow, a first “naïve” approach



Deep dive

Getting video transcripts

Konami'ing your way into video transcripts



<https://github.com/Kakulukian/youtube-transcript>

```
import { storeTranscript } from "adapters/transcript-repository";
import { YoutubeTranscript } from "youtube-transcript";

export const handler = async (event: {
  youtubeVideoUrl: string;
  requestId: string;
}) => {
  const { youtubeVideoUrl, requestId } = event;
  const transcript = await YoutubeTranscript.fetchTranscript(youtubeVideoUrl);
  const sentences = Array.from(getSentencesFromYoutubeTranscript(transcript));

  await storeTranscript(requestId, sentences.join("\n"));
};

function* getSentencesFromYoutubeTranscript(transcript: { text: string }[]) {
  let currentSentence: string[] = [];
  let i = 0;
  do {
    const { text } = transcript[i];

    currentSentence.push(text);

    if (text.endsWith(".")) {
      yield currentSentence.join(" ").replaceAll("\n", " ");
      currentSentence = [];
    }
    i++;
  } while (i < transcript.length);

  yield currentSentence.join(" ").replaceAll("\n", " ");
}
```

~ Fast, as it retrieves already generated transcripts from youtube

~ ...but it's brittle, based on a non official youtube API, does not work on some languages

Deep dive

Prompting techniques with Claude

You are a video transcript summarizer.
Here is a video transcription that you must summarize:

```
<transcript>{{transcript}}</transcript>
```

Summarize this transcript in a third person point of view in 10 sentences.
Identify the speakers and the main topics of the transcript and add them in the output as well.

Do not add or invent speaker names if you not able to identify them.

You must output the summary JSON format conforming to this JSON schema:

```
{
  "type": "object",
  "properties": {
    "speakers": {
      "type": "array",
      "items": {
        "type": "string"
      }
    },
    "topics": {
      "type": "string"
    },
    "summary": {
      "type": "array",
      "items": {
        "type": "string"
      }
    }
  }
}
```

- Clear without ambiguity, no room for assumptions
- Don't be polite, be straightforward
- Concrete example/ schema for structured outputs



Deep dive

Putting words in Claude's mouth

For structured output, we'll need to give hints to the Assistant by prefilling its response

With Claude's Completion API
(Legacy)

```
\n\nHuman:{{prompt}}\n\nAssistant:{
```

Message API

```
{  
  "role": "user",  
  "content": "{{Prompt}}"  
},  
{  
  "role": "assistant",  
  "content": "{"
```

AI



Deep dive

Putting words in Claude's mouth

```
Human:generate a random json


Assistant:
Here is a randomly generated JSON object:

```json
{
 "users": [
 {
 "id": "3c44fb2a",
 "name": "John Smith",
 "age": 37,
 "address": {
 "street": "123 Main St",
 "city": "Anytown",
 "state": "CA"
 }
 }
],
 {
```



```
Human:generate a random json

Assistant:{
 "users": [
 {
 "id": "3c44fb2a",
 "name": "John Smith",
 "age": 37,
 "address": {
 "street": "123 Main St",
 "city": "Anytown",
 "state": "CA"
 }
 },
 {
 "id": "7b349571",
 "name": "Jane Doe",
 "age": 29.
```



AI



# Deep dive

## Invoking the model - state machine definition

```
new LambdaInvoke(this, "generate-model-parameters", {
 lambdaFunction: generateModelParameters,
 payload: TaskInput.fromObject({
 "requestId.$": "$$.Execution.Name",
 }),
}).addCatch(failState)
}

.next(
 new CustomState(this, "bedrock-invoke-model", {
 stateJson: {
 Type: "Task",
 Resource: "arn:aws:states:::bedrock:invokeModel",
 Parameters: {
 ModelId: "anthropic.claude-v2:1",
 Input: {
 "S3Uri.$": `${$.Payload.modelParameters}`,
 },
 ContentType: "application/json",
 },
 },
 ResultSelector: {
 "requestId.$": "$$.Execution.Name",
 "summaryTaskResult.$":
 "States.StringToJson(States.Format('{{{}}}', $.Body.completion))",
 },
 }),
).addCatch(failState)
.next(
 new LambdaInvoke(this, "generate-audio-from-summary", {
 lambdaFunction: generateAudioFromSummary,
 }).addCatch(failState)
)
```

1- Generate model parameters and write to an s3 bucket

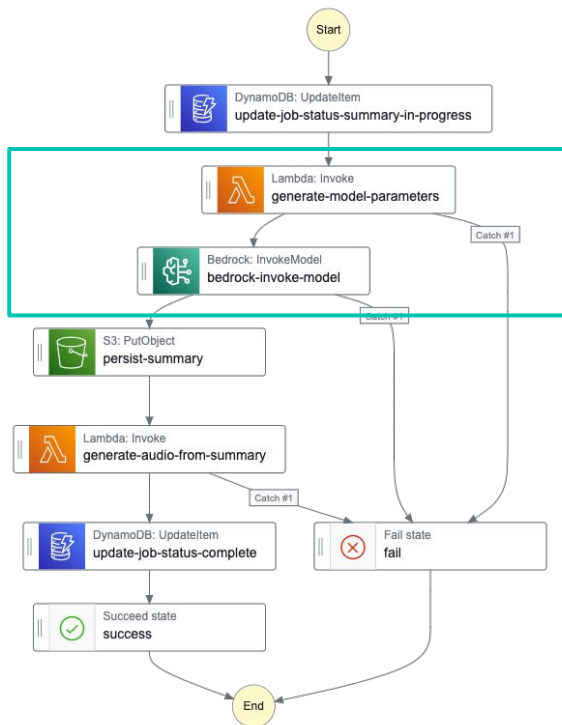
2- Direct invoke to claude model

3- Making sure that JSON is well formed



# Deep diving

## Invoking the model



```
{
 "name": "bedrock-invoke-model",
 "output": {
 "jobId": "37e24860-efdb-4228-a407-63163a2ec051",
 "summaryTaskResult": {
 "speakers": [
 "Michael Hildebrandt",
 "Andre",
 "Fad Sidiki"
],
 "topics": "Making semantic search and AI real, building a production-ready app with semantic sea",
 "summary": [
 "Michael, Andre and Fad talked about making semantic search and AI real to build production-re",
 "They discussed how search is evolving from just keywords to conversational, semantic search.",
 "Customers want search to be more natural, not just entering keywords.",
 "To make semantic search real, companies need to choose appropriate AI models, build semantic",
 "Andre discussed Amazon Bedrock which provides foundation models to build AI applications.",
 "Companies can customize models using techniques like fine-tuning and retrieval-augmented gene",
 "Elasticsearch provides semantic search capabilities to retrieve relevant data and context.",
 "Michael explained Elasticsearch's vector search capabilities for storing, indexing and search",
 "He highlighted performance optimizations done in Elasticsearch like parallel segment search.",
 "Fad discussed an ecommerce use case of enriching product catalog data to improve search recal",
 "He explained using lexical expansion with fine-tuned domain-specific language models to expan",
 "This helps match user search queries better than raw product data.",
 "He emphasized the need for a flexible platform like Elasticsearch to run experiments between"
]
 }
 }
},
 "outputDetails": {
```

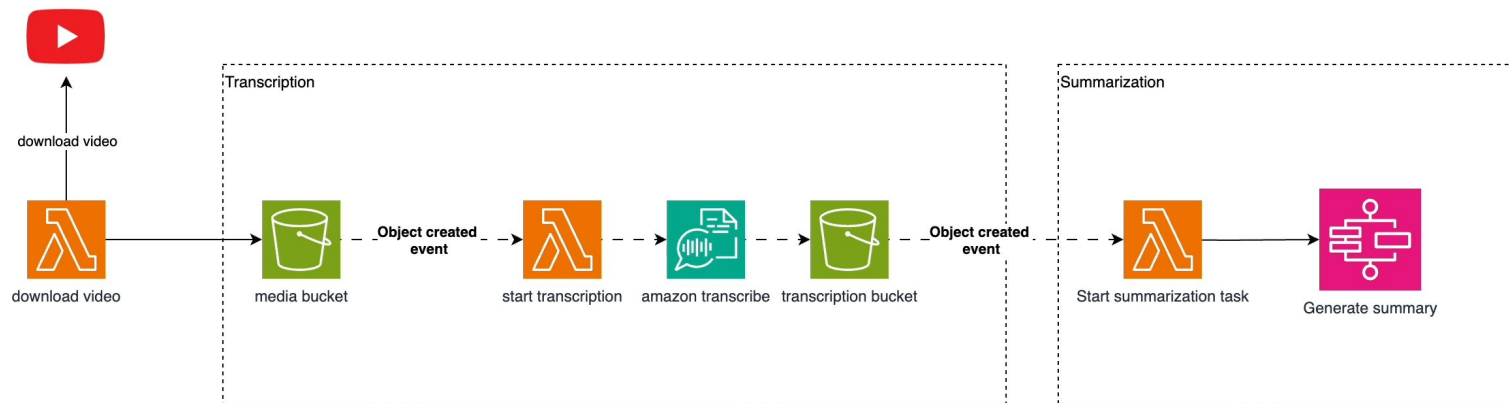




**Can we do better ?**

# A better solution

## Enters Amazon Transcribe



- ~ A more stable solution that supports many languages and more media types
- ~ Some important architectural side effects: The system is now **event based** and **asynchronous**

**This is nice...but how can I scale  
my architecture ?**


# 03

## Scaling the architecture with EDA




# Scaling the architecture

## Orchestration or Choreography ?



~ Orchestration: Greater control over complex workflows. Complex to evolve when collaborating with external contexts

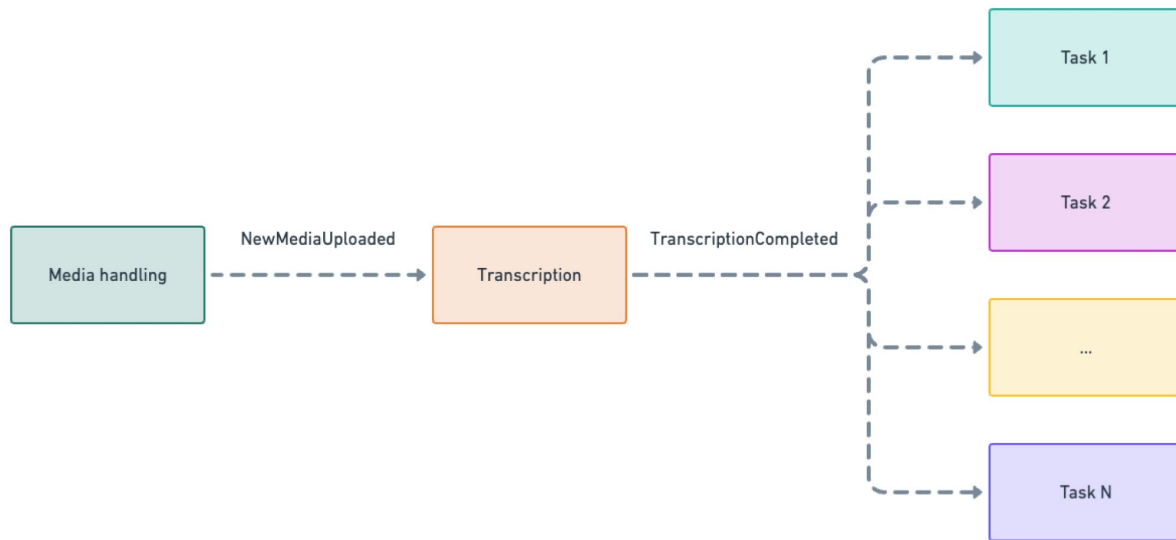


~ Choreography: Flexible and easier to extend but harder to Observe

**Which pattern works best ?**

# Scaling the architecture

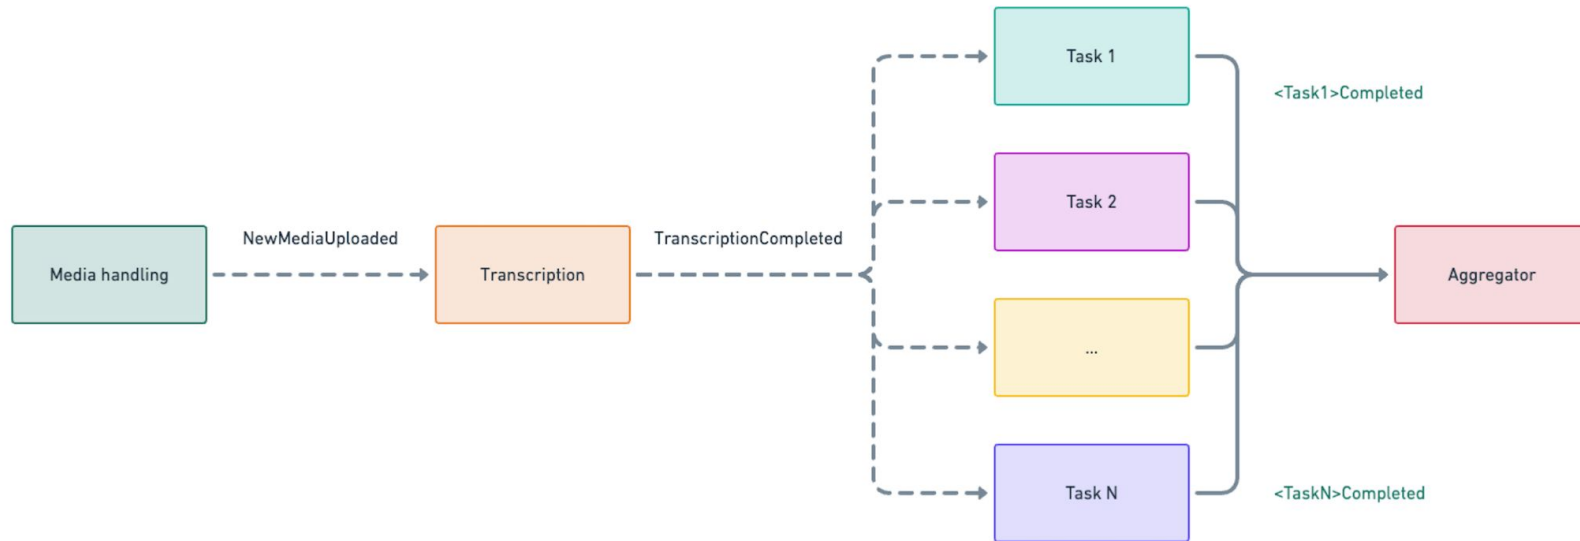
Let's try with Choreography



- ~ Extensible: Tasks can be added or removed without impacting the entire system
- ~ Resilient: If a task fails, it does not impact the entire application

# Scaling the architecture

## Scatter-gather pattern

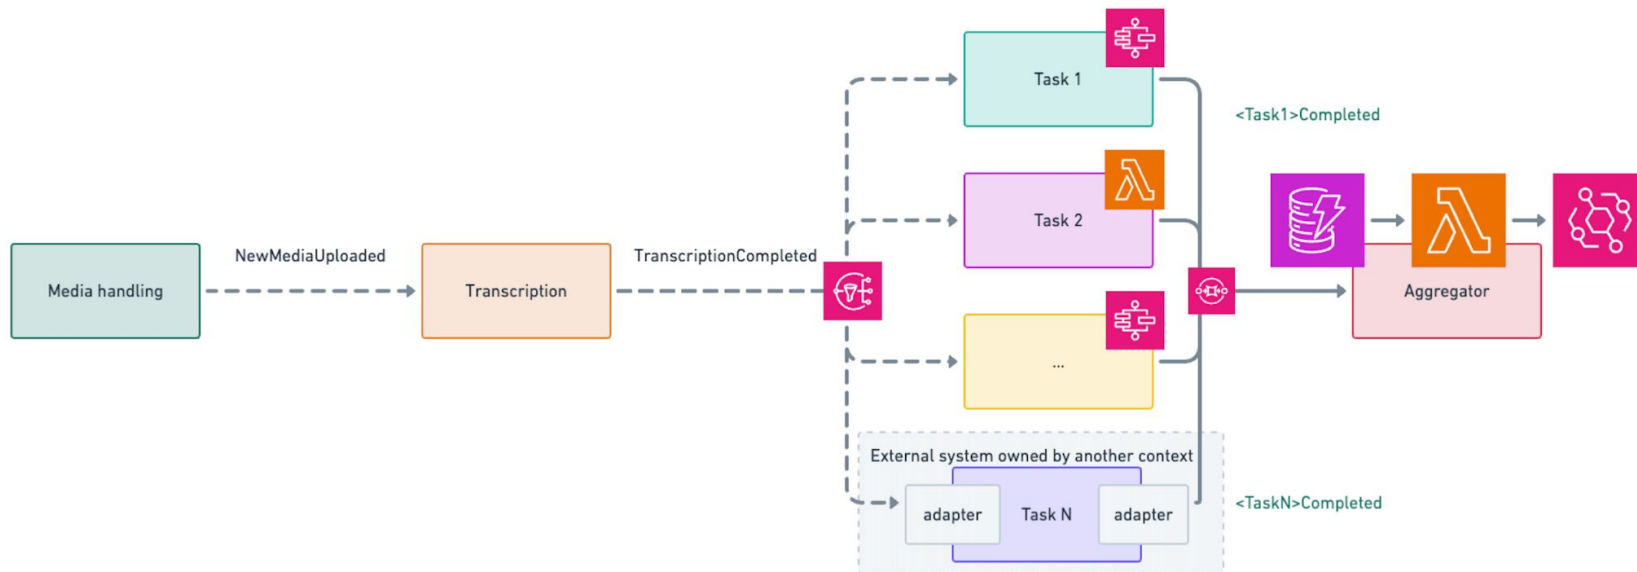


~ Aggregator: A central component handling the events from the scatter phase.



# Scaling the architecture

## Scatter-gather pattern



**...Yes, but**

# Scaling the architecture

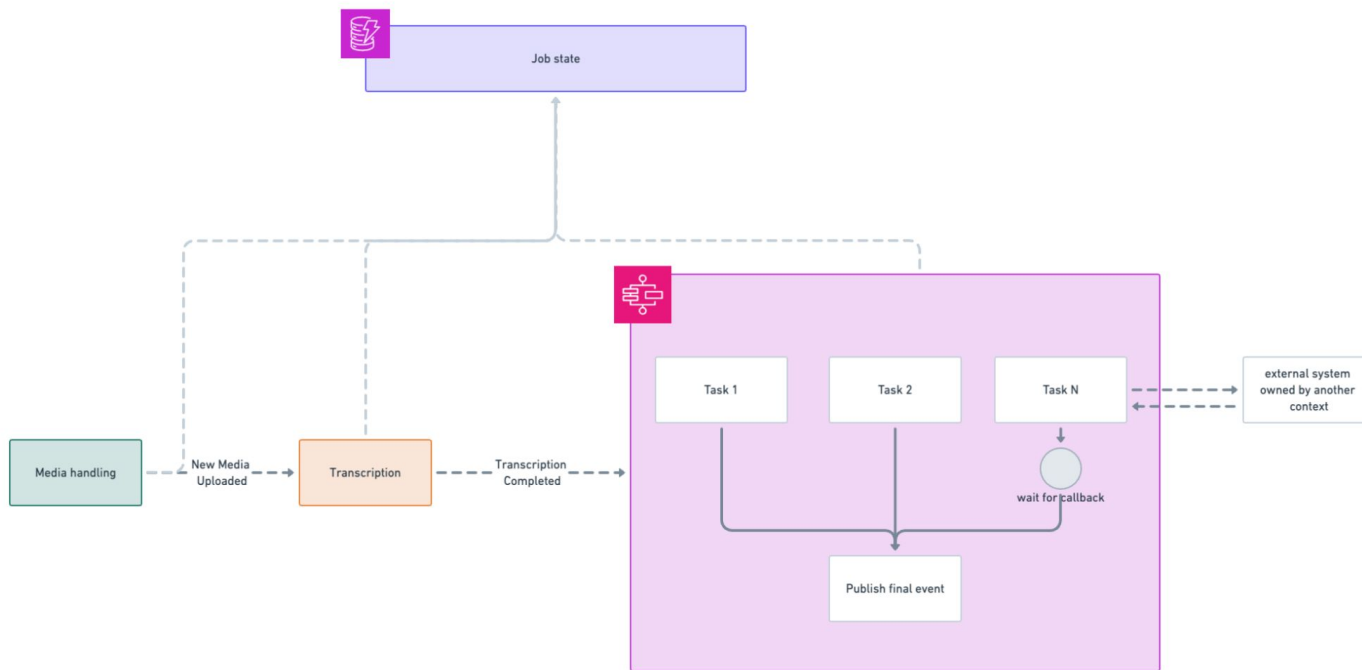
Aggregator pattern can be hard to tame

- What if...
  - a task fails
  - or never publishes the expected event
- How to make sure that failure is handled properly ?
- How to monitor the overall progress of a given Job ?
- How to make sure that my system is observed properly

**So...Orchestration or Choreography ?**

# Scaling the architecture

Orchestration or Choreography...Why not both ?



# 04

## Apify-ing the solution



**APIs, like diamonds, are forever**

*Joshua Bloch*

# Apify-ing the solution

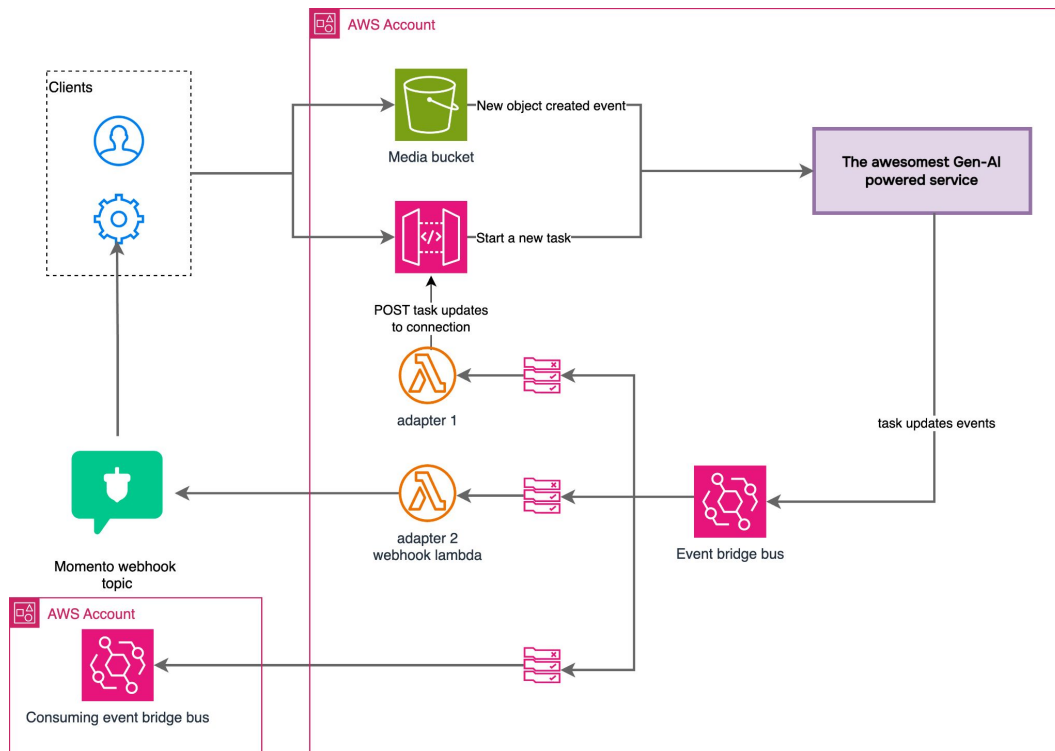
Principles on designing good “public” Event-Driven APIs

- Be **conservative** on what you send (Postel’s Law)
- Be mindful of the **observed behavior** of your APIs (Hyrum’s law)
- Make clear distinction between **your private** and **public events**
- Adopt a **sane** versioning strategy
- Adopt standards: **CloudEvents** & **AsyncAPI**



# Apify-ing the solution


## The awesomest GenAI powered video summarizer



# 05

## Demo





# Thanks !

Questions ?

Source code 🖱️

